

CSci 4061 Programming Assignment 2

Due: Friday, 03/09/2018, 23:59.

Weight: 100 points, 10% of total grade.

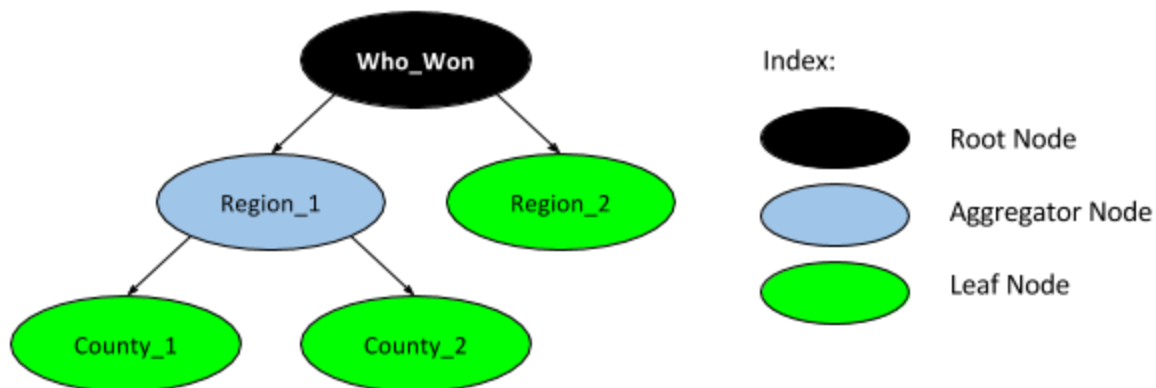
Rule: The assignment must be done in a group of 1 or 2. However, group of 2 is recommended.

Introduction

The purpose of this lab is to learn using the Operating System's I/O and file system calls. In particular, you will do some practices with File I/O operations, File pointers and buffering, File descriptors and I/O redirections, file system, and directories. In the previous assignment, we provided several utility functions, e.g., to read, parse, and write a file. In this assignment, you will need to create your own functions to do those functions.

The Overview of Our Existing Vote Count Application

Recall that our vote count application has three types of node: (1) root node, (2) aggregator node, and (3) leaf node. An example of the application can be seen in the figure below:



In **our existing** vote count application:

- 1) The leaf node will execute a "leafcounter" executable which will read an input file and write out the summary of the input file into an output file.
- 2) The aggregator node will execute a "aggregate_votes" executable which will read n intermediate results and aggregate them into a single output file.
- 3) The root node has two main functions: (1) to parse an <input.txt> which contains the information of the DAG (Directed Acyclic Graph) and (2) to execute "Output_Who_Won" executable which will read n intermediate results, aggregate them, and output the winner to an output file.

Your Assignment

In this assignment, you will create a new breed of the Vote Counter application in which the structure of the graph is given in the form of a directory structure. For example, the figure above can be translated into a structure of 5 directories with a root node of Who_Won directory as follows:

```
Bob@foo:/$ ls
Who_Won

Bob@foo:/$ cd Who_Won; ls
Region_1 Region_2

Bob@foo:/Who_Won$ ls Region_2
votes.txt //Contains the votes for Region_2.

Bob@foo:/Who_Won$ cd Region_1; ls
County_1 County_2

Bob@foo:/Who_Won/Region_1$ ls County_1; ls County_2
votes.txt //Contain the votes for County_1.
votes.txt //Contain the votes for County_2.
```

Your application will need to traverse the directories and aggregate the votes as needed. Depending on your implementation, your application may traverse the directories in a depth-first search ([DFS](#)) or in a breadth-first search ([BFS](#)) manner. Furthermore, you will need to write your own function to read and write the result to a file. In general, your assignment is to create **three** executables, i.e., Leaf_Counter, Aggregator, and Vote_Counter that will be discussed below.

Setup

You may use your assignment 1 as a starting point for the project. However, it is **recommended** to start the assignment from scratch as there will be many changes to your assignment 1 code. However, you may use some functions from your assignment 1 to do this assignment, e.g., the Makeargv.h. However, you are **not** allowed to use any binaries that we have provided for programming assignment 1. There are three executables that you will need to create for this assignment. (**Note:** the name of the executables must be the same with the following):

Executable 1: Leaf_Counter

This executable will act as the leaf nodes.

Input:

This program has **one** argument and will be executed as follows:

- ```
./Leaf_Counter <path>
- <path>: is the relative path to a directory.
```

### Main Functionality:

- This program has a similar functionality with your first programming assignment in which it will read the votes file located at the path, i.e., "<path>/votes.txt", aggregate the votes for each candidate, and output the results into a file that is also located in the path.
- The votes file will always be called "votes.txt" and the output file name will always be the path name with ".txt" appended to it. For example, if the  
<path>="Who\_Won/Region\_1/County\_1", the output file will be  
"Who\_Won/Region\_1/County\_1/County\_1.txt".

### Output:

- If the program is called on a directory that does **not** have votes.txt, it writes to the standard output the following message: "Not a leaf node.\n"
- Otherwise, the program has two outputs: (1) output the counting result into an output file whose name is the last directory name in the path appended with ".txt" and (2) write to the standard output the path to the output file ended with a new line character.
- The output file contains a counting result in the form of:  
"<candidate\_1>:<count\_1>,<candidate\_2>:<count\_2>,...,<candidate\_n>:<count\_n>\n"
- If the output file already exists, it will **replace** it with the new one.

### Example: (Look at figure above for the structure of the DAG)

```
Bob@foo:/$./Leaf_Counter Who_Won/Region_1
Not a leaf node.

Bob@foo:/$ ls Who_Won/Region_1/County_1
votes.txt

Bob@foo:/$ cat Who_Won/Region_1/County_1/votes.txt
A
B
A
C

Bob@foo:/$./Leaf_Counter Who_Won/Region_1/County_1
Who_Won/Region_1/County_1/County_1.txt

Bob@foo:/$ ls Who_Won/Region_1/County_1
vote.txt County_1.txt

Bob@foo:/$ cat Who_Won/Region_1/County_1/County_1.txt
A:2,B:1,C:1
```

## Executable 2: Aggregate\_Votes

This program will act as the aggregator nodes.

### Input:

This program has **one** argument and will be executed as follows:

```
./Aggregate_Votes <path>
```

- <path>: is the relative path to a directory.

### Main Functionality:

- This program has a similar functionality with your first programming assignment in which it will aggregate all the voting results of **all subdirectories**, i.e., its children, that are defined in <path> and write the aggregation result to a file in <path>. The aggregation result file will have the same name as the current node appended with “.txt”
- If we call this program on a leaf node, then it will execute the Leaf\_Counter program.
- Each subdirectory (or child) can either be: (1) a leaf node or (2) a non-leaf node. If the child is a leaf, then this program will spawn a child process to execute the Leaf\_Counter program. If the child is a non-leaf, then this program will spawn a child process to execute another Aggregate\_Votes program.
- The child process **must not** print their result into the standard output.
- If a directory is not a leaf node, however, contains a “votes.txt”, then this program **must ignore** the “votes.txt”.

### Output:

- The program has two outputs: (1) output the aggregation result into an output file located in <path> with a file name of the directory name appended with “.txt” and (2) write to the standard output the output file path ended with a new line character.
- The output file contains a counting result in the form of:  
“<candidate\_1>:<count\_1>,<candidate\_2>:<count\_2>,...,<candidate\_n>:<count\_n>\n”
- If the output file already exists, it will **replace** it with the new one.

**Example:** (Look at figure above for the structure of the DAG)

```
Bob@foo:/$ ls Who_Won/Region_1/County_1
votes.txt County_1.txt

Bob@foo:/$./Aggregate_Votes Who_Won/Region_1/County_1
Who_Won/Region_1/County_1/County_1.txt

Bob@foo:/$ cat Who_Won/Region_1/County_1/County_1.txt
A:2,B:1,C:1

Bob@foo:/$ cat Who_Won/Region_1/County_2/votes.txt
A
D
```

```

Bob@foo/:$./Aggregate_Votes Who_Won/Region_1
Who_Won/Region_1/Region_1.txt

Bob@foo/:$ ls /Who_Won/Region_1/County_2 //Check the output of County_2
votes.txt County_2.txt

Bob@foo/:$ cat /Who_Won/Region_1/Region_1.txt
A:3,B:1,C:1,D:1

```

### Executable 3: Vote\_Counter

This program is the root node of the whole vote counting application.

#### Input:

**Revision 03/06/18:** This program has ~~zero~~ **one** argument and will be executed as follows:

```

./Vote_Counter <root_path>
- <root_path>: the relative path from the current directory to the root.

```

#### Main Functionality:

- The working directory where the program is executed will be the root of the program.
- This program basically will call the Aggregate\_Votes program, read output file of the aggregation result, and **append** the winner into the output file.
- Similarly, child process **should not** write any messages to the standard output.
- This program can also be called from any parts of the DAG.

#### Output:

- The program has one output: write to the standard output the output file path ended with a new line character.
- This program will also append the aggregation result with a winner as follows:  
"Winner: <candidate\_n>\n"

**Example:** (Look at figure above for the structure of the DAG)

```

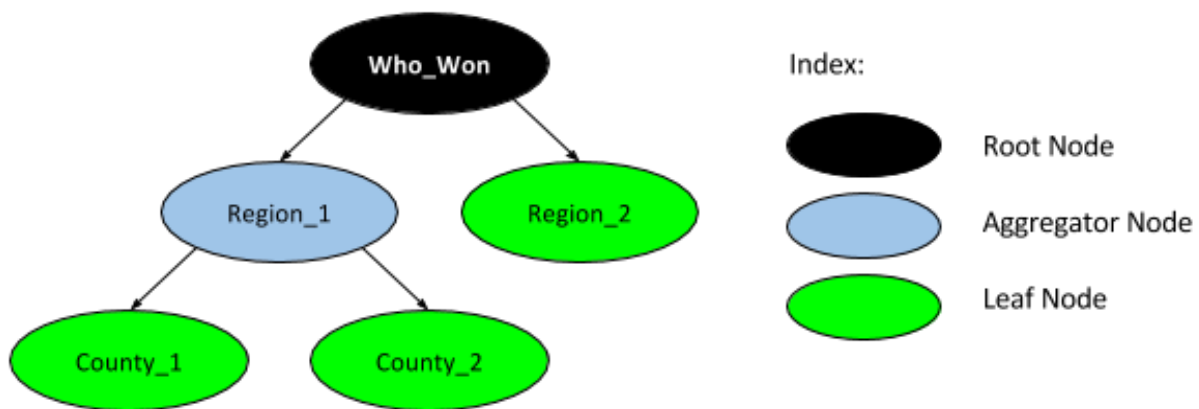
Bob@foo:/$ ls Who_Won
Region_1 Region_2
Bob@foo:/$ cat Who_Won/Region_2/votes.txt
E
E
E
E
Bob@foo:/$./Vote_Counter Who_Won
Who_Won/Who_Won.txt
Bob@foo:/$ cat Who_Won/Who_Won.txt
A:3,B:1,C:1,D:1,E:4 //Written by the Aggregate_Votes on Who_Won
Winner:E //Appended by the Vote_Counter

```

## Putting It Together

- Each executable must be able to be executed on its own without executing the Vote\_Counter program. For example, we should be able to run Leaf\_Counter and Aggregate\_Votes independently.
- Both Aggregate\_Votes and Vote\_Counter can be executed on **any parts** of the DAG. For example, we can run an Aggregate\_Votes or Vote\_Counter on a leaf node.

## Example



- 1) Calling the Leaf\_Counter program on County\_1 will count the data from the leaf node and print the result file name into the standard output.
- 2) Calling the Aggregate\_Votes on Who\_Won will spawn two processes, one for the Region\_2 and one for the Region\_1. For the Region\_2, the Aggregate\_Votes will execute a Leaf\_Counter. For the Region\_1, it will execute another Aggregate\_Votes on Region\_1. Then, the Who\_Won aggregator will read the output of both children (the file names of their output), parse the files and write out the aggregation result to the current directory. Lastly, it will print the aggregation result file name into the standard output.
- 3) Calling the Vote\_Counter on Who\_Won will execute the Aggregate\_Votes on Who\_Won and append the winner into the output\_file.
- 4) Calling an Aggregate\_Votes on County\_1 will execute a Leaf\_Counter program on County\_1 because County\_1 is a leaf node.
- 5) Calling a Vote\_Counter on County\_1 will execute the Aggregate\_Votes program which will execute the Leaf\_Counter program on County\_1. Then, the Vote\_Counter will append the winner to the output file.

# Rules and Hints

## Assignment Rules

- You cannot use the library call "system".
- Every executable must follow the naming convention that is described by the documentation of this assignment, i.e., Leaf\_Counter, Aggregate\_Votes, and Vote\_Counter.
- Both Aggregate\_Votes and Vote\_Counter executables must be able to be executed from any parts of the DAG.
- In contrast to your Programming Assignment 1, **there is no maximum number of children**. Your program must be dynamically able to handle any number of children.
- You **must not** use "cd" and "ls"

## Useful System Calls and Functions

open, close, read, write, fopen, fgets, fclose, fork, execl, opendir, readdir  
Please consult the manpage for the details on how to use them.

## Simplifying Assumption

1. There is always a winner in the voting. (No ties).
2. The maximum size of a file name is 255 bytes.
3. The maximum size for each I/O operation message, e.g., each read or write, is 1024 bytes.

## Error Handling

- You are expected to check the return value of all system calls that you use in your program and check for error conditions. Also, each program should check to make sure the proper number of arguments is used when it is executed. If your program encounters an error, a useful error message should be printed to the screen.
- Your program should be robust; it should try to recover if possible. If the error prevents your program from functioning, then it should exit after printing the error message. (The use of the perror function for printing error messages from library calls is encouraged.)
- Since we are going to redirect the standard output of some programs, you can try to debug your code by printing into the standard error instead.
- **Your program will be intentionally tested with errors.**

## Deliverables

You must submit **single zip file** which contains **at least** the following files:

1. Makefile
2. README
3. At least 3 c files which contain the source code of your Leaf\_Counter, Aggregate\_Votes, and Vote\_Counter

**DO NOT** include your executables. Your submission must be able to be compiled without additional files.

**Revision 02/28/18:** Please zip the files directly (instead of zipping the folder that contains your submission files) so that the c files are located in the root of the zip file.

### Makefile

Your Makefile should be able to be executed as follows:

1. make Leaf\_Counter: creates the Leaf\_Counter executable.
2. make Aggregate\_Votes: creates the Aggregate\_Votes executable.
3. make Vote\_Counter: creates the Vote\_Counter executable.
4. make: creates all three executables.
5. make clean: remove all generated executables.

### README

You must include a README file, named "README" without extensions, which describes your program.

It needs to contain the following:

- The purpose of your program
- Your x500 and the x500 of your partner
- Your lecture section and your partner's lecture section
- Your partner's and your individual contributions
- **Specify whether you are doing the extra credit or not.**

At the top of your README file and main C source file please include the following comment:

```
/*login: x500_1, x500_2
date: mm/dd/yy
name: full_name1, full_name2
id: id_for _first_name, id_for _second_name */
```



## Grading Rubric (Revision 03/06/18)

This the exact grading rubric:

- 10% Following the Deliverable Instructions (NO PARTIAL CREDIT)
  - All source codes, including .c, .h, and Makefile, are located at the **root of the zip** and **not** in a folder that you zip. Please **do not** zip the directory that contains your files, instead, select all your files then zip them directly.
  - All executables follow the correct naming convention
- 5% README and Makefile.
  - The Makefile can be compiled **with GCC 4.9.2**
  - The README is complete
- 15% Documentation within code, coding, and style. (5% each)
- 70% Test Cases.
  - Functionality test cases (55%)
  - Error handling test cases (15%)

## Extra Credit (10 Additional Points)

For the extra credit, you will need to make the Vote\_Counter program to be aware if the graph has a cycle. We will use symbolic link to create a cycle to the graph. In such case, the Vote\_Counter program should work as before, however, it will append an additional information to the output file that specifies where the cycle is located. For example: assume there is a symbolic link from County\_1 to Who\_Won. Then, your Vote\_Counter application should have the following output:

```
Bob@foo:/$./Vote_Counter Who_Won
Who_Won/Who_Won.txt

Bob@foo:/$ cat Who_Won/Who_Won.txt
A:3,B:1,C:1,D:1,E:4 //Written by the Aggregate_Votes on Who_Won
Winner:E //Appended by the Vote_Counter
There is a cycle from County_1 to Who_won. //Appended by the extra credit.
```

Clarification 03/08/18:

- (1) In this extra credit, you will **not** given a “**bad symbolic link**”, i.e., a link that points to some random directories that are not part of the DAG.
- (2) The link is also **guaranteed** to link the node’s direct ancestors, e.g., its parent, grandparent, etc., and will not point to its sibling. For example, there can be a link from County\_1 to Region\_1 or Who\_Won, however, there will **not** be a link from County\_1 to Region\_2.
- (3) You **cannot** assume that there is only one symbolic link. For example, County\_1 directory may have two symbolic links to Region\_1 and Who\_Won.